

## Problem Set 1

**Instructions:** This problem set is not a part of your final grade. Solve at least 3 of these problems and submit your solutions. We will not look at them, but will provide solutions at the end of the deadline for you to compare your answers. At the end of the term, the fact that you turned something in may be used as a part of the process for determining your final grade. More on that process after the midterm.

1. Where we are doing supervised learning, we have mostly assumed a deterministic function. Imagine instead a world where we are trying to capture a non-deterministic function. In this case, we might see training pairs where the  $x$  value appears several times, but with different  $y$  values. For example, we might use attributes of humans to the probability that they have had chicken pox. In that case, we might see the same kind of person many times but only sometimes they may have had chicken pox.

We would like to build a learning algorithm that will compute the probability that a person has chicken pox. So, given a set of training data where each instance is mapped to 1 for *true* or 0 for *false*:

1. Derive the proper error function to use for finding the ML hypothesis using Bayes' Rule. You should go through a similar process as the one used to derive least squared error in the lessons.
  2. Compare and contrast your result to the rule we derived for a deterministic function perturbed by zero-mean gaussian noise. What would a normal neural network using sum of squared errors do with these data? What if the data consisted of  $x,y$  pairs where  $y$  was an estimate of the probability instead of 0s and 1s?
2. Design a two-input perceptron that implements the boolean function  $A \wedge \neg B$ . Design a two-layer network of perceptrons that implements  $A \oplus B$  (where  $\oplus$  is XOR).
  3. Derive the perceptron training rule and gradient descent training rule for a single unit with output  $o$ , where  $o = w_0 + w_1x_1 + w_1x_1^2 + \dots + w_nx_n + w_nx_n^2$ . What are the advantages of using gradient descent training rule for training neural networks over the perceptron training rule?
  4. Explain how one can use Decision Trees to perform regression? Show that when the error function is squared error that the expected value at any leaf is the mean. Take the Boston Housing dataset (<http://lib.stat.cmu.edu/datasets/boston>) and use Decision Trees to perform regression.
  5. Suggest a lazy version of the eager decision tree learning algorithm ID3. What are the advantages and disadvantages of your lazy algorithm compared to the original eager algorithm?
  6. Imagine you had a learning problem with an instance space of points on the plane and a target function that you knew took the form of a line on the plane where all points on one side of the line are positive and all those on the other are negative. If you were constrained to only use decision tree or nearest-neighbor learning, which would you use? Why?
  7. Give the VC dimension of these hypothesis spaces, briefly explaining your answers:
    1. An origin-centered circle (2D)
    2. An origin-centered sphere (3D)